

# Proj\_emp\_Turn\_analysis

April 14, 2024

## Employee Turnover Analytics

### Project Statement:

Portobello Tech is an app innovator that has devised an intelligent way of predicting employee turnover within the company. It periodically evaluates employees' work details including the number of projects they worked upon, average monthly working hours, time spent in the company, promotions in the last 5 years, and salary level.

Data from prior evaluations show the employee's satisfaction at the workplace. The data could be used to identify patterns in work style and their interest to continue to work in the company.

The HR Department owns the data and uses it to predict employee turnover. Employee turnover refers to the total number of workers who leave a company over a certain time period.

### Importing libraries

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

### Loading excel file

```
[2]: data = pd.read_excel('1673873196_hr_comma_sep.xlsx')
```

```
[3]: data
```

```
[3]:
```

	satisfaction_level	last_evaluation	number_project	\
0	0.38	0.53	2	
1	0.80	0.86	5	
2	0.11	0.88	7	
3	0.72	0.87	5	
4	0.37	0.52	2	
...	...	...	...	
14994	0.40	0.57	2	
14995	0.37	0.48	2	
14996	0.37	0.53	2	

14997	0.11	0.96	6
14998	0.37	0.52	2

	average_monthly_hours	time_spend_company	Work_accident	left	\
0	157	3	0	1	
1	262	6	0	1	
2	272	4	0	1	
3	223	5	0	1	
4	159	3	0	1	
...	...	...	...	...	
14994	151	3	0	1	
14995	160	3	0	1	
14996	143	3	0	1	
14997	280	4	0	1	
14998	158	3	0	1	

	promotion_last_5years	sales	salary
0	0	sales	low
1	0	sales	medium
2	0	sales	medium
3	0	sales	low
4	0	sales	low
...	...	...	...
14994	0	support	low
14995	0	support	low
14996	0	support	low
14997	0	support	low
14998	0	support	low

[14999 rows x 10 columns]

## Perform steps

1. Perform data quality check by checking for missing values if any.

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   satisfaction_level     14999 non-null  float64
1   last_evaluation       14999 non-null  float64
2   number_project        14999 non-null  int64
3   average_monthly_hours 14999 non-null  int64
4   time_spend_company    14999 non-null  int64
5   Work_accident         14999 non-null  int64
```

```

6   left                14999 non-null   int64
7   promotion_last_5years 14999 non-null   int64
8   sales                14999 non-null   object
9   salary               14999 non-null   object
dtypes: float64(2), int64(6), object(2)
memory usage: 1.1+ MB

```

```
[5]: data.isnull().sum()
```

```

[5]: satisfaction_level    0
    last_evaluation        0
    number_project         0
    average_monthly_hours  0
    time_spend_company     0
    Work_accident          0
    left                   0
    promotion_last_5years  0
    sales                  0
    salary                 0
dtype: int64

```

- No missing values in the dataset

2. Understand what factors contributed most to employee turnover by EDA. 2.1. Draw a heatmap of the Correlation Matrix between all numerical features/columns in the data. 2.2. Draw the distribution plot of Employee Satisfaction (use column satisfaction\_level) Employee Evaluation (use column last\_evaluation) Employee Average Monthly Hours (use column average\_monthly\_hours) 2.3. Draw the bar plot of Employee Project Count of both employees who left and who stayed in the organization (use column number\_project and hue column left) and give your inferences from the plot.

```
[6]: # Heatmap of the Correlation Matrix
```

```

[7]: plt.figure(figsize=(10,8))
    sns.heatmap(data.corr(),annot=True)
    plt.show()

```



```
[14]: # Distribution Plot
      #Employee Satisfaction
      # Employee Evaluation
      # Employee Average Monthly Hours
```

```
[12]: plt.figure(figsize=(12,4))


plt.subplot(1,3,1)
sns.distplot(data['satisfaction_level'],kde=True,color='purple')
plt.title('Emp Satisfaction level')

plt.subplot(1,3,2)
sns.distplot(data['last_evaluation'],kde=True,color='red')
plt.title('Emp evaluation level')

plt.subplot(1,3,3)
```

```
sns.distplot(data['average_monthly_hours'],kde=True,color='green')
plt.title('Emp Monthly hours')

plt.tight_layout()
plt.show()
```



- For more clear visualization, I am using Histogram plot

[13]: *# Histplot for data visualization*


```
plt.figure(figsize=(12,4))

plt.subplot(1,3,1)
sns.histplot(data['satisfaction_level'],kde=True,color='purple')
plt.title('Emp Satisfaction level')

plt.subplot(1,3,2)
sns.histplot(data['last_evaluation'],kde=True,color='red')
plt.title('Emp evaluation level')

plt.subplot(1,3,3)
sns.histplot(data['average_monthly_hours'],kde=True,color='green')
plt.title('Emp Monthly hours')


plt.tight_layout()
plt.show()
```



```
[16]: # Draw the bar plot of Employee Project Count of both employees who left and
      ↪ who stayed in the organization
      # (use column number_project and hue column left) and give your inferences
      ↪ from the plot.
```


```
[21]: # Employee who left and stayed based on Number of Project

sns.countplot(x=data['number_project'],hue='left',data=data)
plt.show()
```




```
[18]: # Employee who left and stayed based on Salary

sns.countplot(x=data['salary'],hue='left',data=data)
plt.show()
```



```
[23]: # Employee who left and stayed based on Sales

sns.countplot(x=data['sales'],hue='left',data=data)
plt.xticks(rotation=90)
plt.show()
```



3. Perform clustering of Employees who left based on their satisfaction and evaluation.
  - 3.1. Choose columns satisfaction\_level, last\_evaluation and left.
  - 3.2. Do KMeans clustering of employees who left the company into 3 clusters.
  - 3.3. Based on the satisfaction and evaluation factors, give your thoughts on the employee clusters.

```
[24]: data.columns
```

```
[24]: Index(['satisfaction_level', 'last_evaluation', 'number_project',
          'average_monthly_hours', 'time_spend_company', 'Work_accident', 'left',
          'promotion_last_5years', 'sales', 'salary'],
          dtype='object')
```

```
[25]: # Choose columns satisfaction_level, last_evaluation and left
```

```
[26]: cluster_data = data[['satisfaction_level', 'last_evaluation', 'left']]
      left_emp_data = cluster_data[cluster_data['left'] == 1]
```



```
left_emp_data.head(10)
```

```
[26]:
```

	satisfaction_level	last_evaluation	left
0	0.38	0.53	1
1	0.80	0.86	1
2	0.11	0.88	1
3	0.72	0.87	1
4	0.37	0.52	1
5	0.41	0.50	1
6	0.10	0.77	1
7	0.92	0.85	1
8	0.89	1.00	1
9	0.42	0.53	1

```
[27]: left_emp_data.shape
      # checking the total number of rows and cols in left_emp_data
```

```
[27]: (3571, 3)
```

```
[35]: # Do KMeans clustering of employees who left the company into 3 clusters.
```

```
[28]: from sklearn.cluster import KMeans
```


```
[29]: kmeans=KMeans(n_clusters=3,random_state=42)
      kmeans.fit(left_emp_data)
```

```
[29]: KMeans(n_clusters=3, random_state=42)
```

```
[30]: left_emp_data['cluster']=kmeans.labels_
```

```
[37]: # Based on the satisfaction and evaluation factors, give your thoughts on the
      ↪ employee clusters.

      plt.figure(figsize=(10,6))
      sns.
          ↪ scatterplot(x='satisfaction_level',y='last_evaluation',hue='cluster',palette='deep',
          ↪ data=left_emp_data)
      plt.show()
```



- Employee left based on High Evaluation and Low Satisfaction level
- Employee left based on Low Evaluation and Average Satisfaction level
- Employee left based on High Evaluation and High Satisfaction level

```
[34]: left_emp_data['cluster'].value_counts()
      # To show the the number of Employee left based on cluster
```

```
[34]: 0    1650
      1     977
      2     944
      Name: cluster, dtype: int64
```

4. Handle the left Class Imbalance using SMOTE technique 4.1.Pre-Process the data by converting categorical columns to numerical columns by Separating categorical variables and numeric variables. Applying `get_dummies()` to the categorical variables. Combining categorical variables and numeric variables. 4.2.Do the stratified split of the dataset to train and test in the ratio 80:20 with `random_state=123`. 4.3.Upsample the train dataset using SMOTE technique from the `imblearn` module.

```
[38]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 10 columns):
 #   Column              Non-Null Count  Dtype
---  -

```

```

0  satisfaction_level      14999 non-null float64
1  last_evaluation        14999 non-null float64
2  number_project         14999 non-null int64
3  average_monthly_hours  14999 non-null int64
4  time_spend_company     14999 non-null int64
5  Work_accident          14999 non-null int64
6  left                  14999 non-null int64
7  promotion_last_5years  14999 non-null int64
8  sales                  14999 non-null object
9  salary                 14999 non-null object
dtypes: float64(2), int64(6), object(2)
memory usage: 1.1+ MB

```

```
[41]: # Separating categorical variables and numeric variables
```

```

df_numerical = data.select_dtypes(include=['int64','float64'])
df_categorical = data.select_dtypes(include=['object'])

df_categorical.head(10)

```

```

[41]:  sales  salary
0  sales    low
1  sales  medium
2  sales  medium
3  sales    low
4  sales    low
5  sales    low
6  sales    low
7  sales    low
8  sales    low
9  sales    low

```

```
[42]: # Applying get_dummies() to the categorical variables.
```

```
df_converted=pd.get_dummies(data=df_categorical)
```

```
[43]: df_converted.head(10)
```

```

[43]:  sales_IT  sales_RandD  sales_accounting  sales_hr  sales_management  \
0         0           0           0           0           0
1         0           0           0           0           0
2         0           0           0           0           0
3         0           0           0           0           0
4         0           0           0           0           0
5         0           0           0           0           0
6         0           0           0           0           0
7         0           0           0           0           0

```

8	0	0	0	0	0
9	0	0	0	0	0

	sales_marketing	sales_product_mng	sales_sales	sales_support	\
0	0	0	1	0	
1	0	0	1	0	
2	0	0	1	0	
3	0	0	1	0	
4	0	0	1	0	
5	0	0	1	0	
6	0	0	1	0	
7	0	0	1	0	
8	0	0	1	0	
9	0	0	1	0	

	sales_technical	salary_high	salary_low	salary_medium
0	0	0	1	0
1	0	0	0	1
2	0	0	0	1
3	0	0	1	0
4	0	0	1	0
5	0	0	1	0
6	0	0	1	0
7	0	0	1	0
8	0	0	1	0
9	0	0	1	0

[44]: *# Combining categorical variables and numeric variables*

```
df_new=pd.concat([df_numerical,df_converted],axis=1)
df_new.head(10)
```

[44]:	satisfaction_level	last_evaluation	number_project	average_monthly_hours	\
0	0.38	0.53	2	157	
1	0.80	0.86	5	262	
2	0.11	0.88	7	272	
3	0.72	0.87	5	223	
4	0.37	0.52	2	159	
5	0.41	0.50	2	153	
6	0.10	0.77	6	247	
7	0.92	0.85	5	259	
8	0.89	1.00	5	224	
9	0.42	0.53	2	142	

	time_spend_company	Work_accident	left	promotion_last_5years	sales_IT	\
0	3	0	1	0	0	
1	6	0	1	0	0	

2	4	0	1	0	0
3	5	0	1	0	0
4	3	0	1	0	0
5	3	0	1	0	0
6	4	0	1	0	0
7	5	0	1	0	0
8	5	0	1	0	0
9	3	0	1	0	0

	sales_RandD	...	sales_hr	sales_management	sales_marketing	\
0	0	...	0	0	0	
1	0	...	0	0	0	
2	0	...	0	0	0	
3	0	...	0	0	0	
4	0	...	0	0	0	
5	0	...	0	0	0	
6	0	...	0	0	0	
7	0	...	0	0	0	
8	0	...	0	0	0	
9	0	...	0	0	0	

	sales_product_mng	sales_sales	sales_support	sales_technical	\
0	0	1	0	0	
1	0	1	0	0	
2	0	1	0	0	
3	0	1	0	0	
4	0	1	0	0	
5	0	1	0	0	
6	0	1	0	0	
7	0	1	0	0	
8	0	1	0	0	
9	0	1	0	0	

	salary_high	salary_low	salary_medium
0	0	1	0
1	0	0	1
2	0	0	1
3	0	1	0
4	0	1	0
5	0	1	0
6	0	1	0
7	0	1	0
8	0	1	0
9	0	1	0

[10 rows x 21 columns]

```
[45]: # Handling class Imbalance technique
```

```
df_new['left'].value_counts()
```

```
[45]: 0    11428  
      1     3571  
      Name: left, dtype: int64
```

```
[46]: X = df_new.drop('left',axis=1)  
      Y = df_new['left']
```

```
[48]: # Do the stratified split of the dataset to train and test in the ratio 80:20  
      ↪with random_state=123.
```

```
from sklearn.model_selection import train_test_split
```

```
[49]: X_train,X_test,Y_train,Y_test = train_test_split(X,Y,train_size=0.8,  
      ↪random_state=123)
```

```
[58]: # Upsample the train dataset using SMOTE technique from the imblearn module.
```

```
from imblearn.over_sampling import SMOTE
```

```
sm = SMOTE(random_state=2)  
X_train_resample,Y_train_resample = sm.fit_resample(X_train,Y_train)
```

```
[57]: Y_train_resample.value_counts()  
      # checking the final values of Y_training set
```

```
[57]: 0     9137  
      1     9137  
      Name: left, dtype: int64
```

5. Perform 5-Fold cross-validation model training and evaluate performance. 5.1. Train a Logistic Regression model and apply a 5-Fold CV and plot the classification report. 5.2. Train a Random Forest Classifier model and apply the 5-Fold CV and plot the classification report. 5.3. Train a Gradient Boosting Classifier model and apply the 5-Fold CV and plot the classification report.
6. Identify the best model and justify the evaluation metrics used. 6.1. Find the ROC/AUC for each model and plot the ROC curve. 6.2. Find the confusion matrix for each of the models. 6.3. From the confusion matrix, explain which metric needs to be used- Recall or Precision?

```
[66]: from sklearn.model_selection import cross_val_score  
      from sklearn.linear_model import LogisticRegression  
      from sklearn.metrics import accuracy_score,roc_auc_score,classification_report,  
      ↪confusion_matrix
```

```
[105]: # Train a Logistic Regression model and apply a 5-Fold CV and plot the
      ↪classification report, roc auc and cross-validation score
```

```
log_reg = LogisticRegression()
log_reg.fit(X_train_resample,Y_train_resample)
```

```
[105]: LogisticRegression()
```

```
[106]: y_pred1 = log_reg.predict(X_test)
```

```
[64]: print('Accuracy score: ',accuracy_score(Y_test,y_pred1))
```

Accuracy score: 0.7616666666666667

- Here, accuracy score is 76.16%

```
[65]: print(classification_report(Y_test,y_pred1))
```

	precision	recall	f1-score	support
0	0.91	0.76	0.83	2291
1	0.50	0.77	0.60	709
accuracy			0.76	3000
macro avg	0.71	0.76	0.72	3000
weighted avg	0.81	0.76	0.78	3000

```
[69]: print(confusion_matrix(Y_test, y_pred1))
```

```
[[1742  549]
 [ 166  543]]
```

```
[71]: print(cross_val_score(log_reg,X_train_resample,Y_train_resample).mean())
```

0.7988404065181272

```
[72]: print(roc_auc_score(Y_test,y_pred1))
```


0.7631170355084193

```
[73]: from sklearn import metrics
      fpr, tpr, threshold = metrics.roc_curve(Y_test, y_pred1)
      print(fpr)
      print(tpr)
      print(threshold)
      roc_auc = metrics.auc(fpr, tpr)
```

```
[0.      0.23963335 1.      ]
[0.      0.76586742 1.      ]
[2 1 0]
```

```
[80]: # Plot ROC AUC graph for Logistic Regression
```

```
plt.title('Logistic Regression ROC AUC graph')
plt.plot(fpr, tpr, 'r', label = 'AUC')
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1])
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



```
[82]: # Train a Random Forest Classifier model and apply a 5-Fold CV and plot the
      ↪ classification report, roc auc and cross-validation score
```

```
from sklearn.ensemble import RandomForestClassifier
```



```
[83]: random_forest=RandomForestClassifier(max_depth=5)
```

```
[103]: random_forest.fit(X_train_resample,Y_train_resample)
```

```
[103]: RandomForestClassifier(max_depth=5)
```

```
[104]: y_pred = random_forest.predict(X_test)
```

```
[87]: print('Accuracy score',accuracy_score(Y_test,y_pred))
```

Accuracy score 0.9573333333333334

```
[96]: print(cross_val_score(random_forest,X_train_resample,Y_train_resample).mean())
```

0.9477399877352705

```
[95]: print(confusion_matrix(Y_test, y_pred))
```

```
[[2218  73]
 [ 55 654]]
```

```
[88]: print(classification_report(Y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.98	0.97	0.97	2291
1	0.90	0.92	0.91	709
accuracy			0.96	3000
macro avg	0.94	0.95	0.94	3000
weighted avg	0.96	0.96	0.96	3000

```
[90]: print(roc_auc_score(Y_test,y_pred))
```


0.9452810685585774

```
[91]: fpr, tpr, threshold = metrics.roc_curve(Y_test, y_pred)
print(fpr)
print(tpr)
print(threshold)
roc_auc = metrics.auc(fpr, tpr)
```

```
[0.          0.03186381 1.          ]
[0.          0.92242595 1.          ]
[2 1 0]
```

```
[94]: # Plot ROC AUC graph for Random Forest Classifier
```

```
plt.title('Random Forest Classifier')
plt.plot(fpr, tpr, 'r', label = 'AUC')
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1])
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



```
[97]: # Train a Gradient Boosting Classifier model and apply a 5-Fold CV and plot the
      ↪ classification report, roc auc and cross-validation score
```

```
from sklearn.ensemble import GradientBoostingClassifier
```

```
[99]: gradient_boost = GradientBoostingClassifier()
```

```
[101]: gradient_boost.fit(X_train_resample,Y_train_resample)
```

```
[101]: GradientBoostingClassifier()
```

```
[102]: y_pred2 = gradient_boost.predict(X_test)
```

```
[107]: print('Accuracy score',accuracy_score(Y_test,y_pred2))
```

Accuracy score 0.964

```
[108]: print(classification_report(Y_test,y_pred2))
```

	precision	recall	f1-score	support
0	0.98	0.97	0.98	2291
1	0.91	0.94	0.92	709
accuracy			0.96	3000
macro avg	0.95	0.96	0.95	3000
weighted avg	0.96	0.96	0.96	3000

```
[109]: print(roc_auc_score(Y_test,y_pred2))
```

0.9550026811235971

```
[110]: print(confusion_matrix(Y_test, y_pred2))
```

```
[[2227  64]
 [  44 665]]
```

```
[111]: print(cross_val_score(gradient_boost,X_train_resample,Y_train_resample).mean())
```

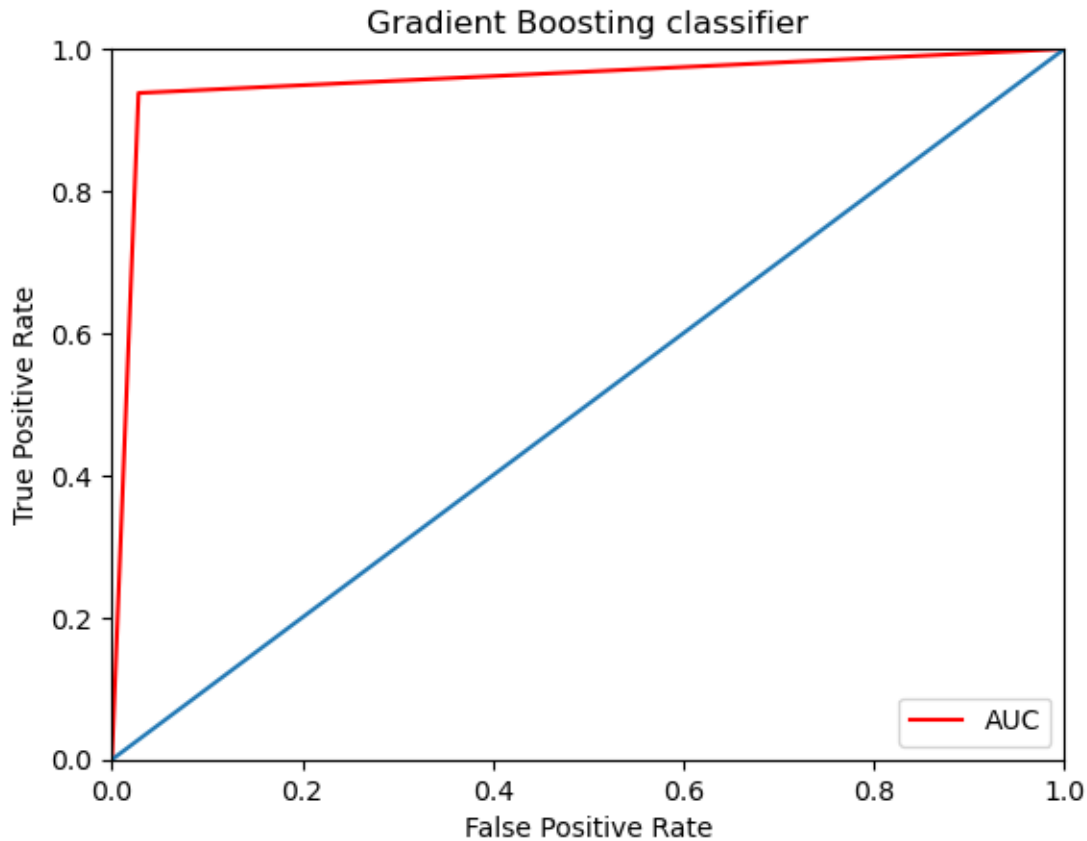
0.9587938933926953

```
[112]: fpr, tpr, threshold = metrics.roc_curve(Y_test, y_pred2)
print(fpr)
print(tpr)
print(threshold)
roc_auc = metrics.auc(fpr, tpr)
```

```
[0.          0.0279354 1.          ]
[0.          0.93794076 1.          ]
[2 1 0]
```

```
[113]: plt.title('Gradient Boosting classifier')
plt.plot(fpr, tpr, 'r', label = 'AUC')
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1])
plt.xlim([0, 1])
```

```
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



- Here, the best fit models were Random Forest and Gradient Boosting classifiers
7. Suggest various retention strategies for targeted employees. 7.1. Using the best model, predict the probability of employee turnover in the test data. 7.2. Based on the below probability score range, categorize the employees into four zones and suggest your thoughts on the retention strategies for each zone. Safe Zone (Green) (Score < 20%) Low Risk Zone (Yellow) (20% < Score < 60%) Medium Risk Zone (Orange) (60% < Score < 90%) High Risk Zone (Red) (Score > 90%).

```
[114]: # Using XGBoost boosting classifier model
```

```
import xgboost as xgb
model = xgb.XGBClassifier()
```

```
[116]: model.fit(X_train_resample, Y_train_resample)
```

```
[116]: XGBClassifier(base_score=None, booster=None, callbacks=None,
                    colsample_bylevel=None, colsample_bynode=None,
                    colsample_bytree=None, device=None, early_stopping_rounds=None,
                    enable_categorical=False, eval_metric=None, feature_types=None,
                    gamma=None, grow_policy=None, importance_type=None,
                    interaction_constraints=None, learning_rate=None, max_bin=None,
                    max_cat_threshold=None, max_cat_to_onehot=None,
                    max_delta_step=None, max_depth=None, max_leaves=None,
                    min_child_weight=None, missing=nan, monotone_constraints=None,
                    multi_strategy=None, n_estimators=None, n_jobs=None,
                    num_parallel_tree=None, random_state=None, ...)
```

```
[117]: y_pred3 = model.predict(X_test)
```

```
[118]: print('Accuracy score', accuracy_score(Y_test, y_pred3))
```

Accuracy score 0.9826666666666667

```
[119]: predicted_prob = model.predict_proba(X_test)
```

```
[120]: predicted_prob[:,1]
```

```
[120]: array([1.0655432e-03, 4.4659537e-04, 4.6851885e-05, ..., 9.8639083e-01,
          4.3727891e-03, 5.5505626e-04], dtype=float32)
```

```
[126]: # Using the best model, predict the probability of employee turnover in the
      ↪ test data
```

```
X_test['leave']=y_pred3
X_test
```

```
[126]:
```

	satisfaction_level	last_evaluation	number_project	\
6958	0.54	0.67	3	
7534	0.72	0.52	3	
2975	0.95	0.61	3	
3903	0.78	0.79	3	
8437	0.60	0.40	3	
...	...	...	...	
1229	0.42	0.55	2	
10593	0.61	0.67	4	
12248	0.87	0.91	4	
3147	0.49	0.71	3	
6623	0.52	0.66	5	

	average_monthly_hours	time_spent_company	Work_accident	\
6958	154	2	0	
7534	143	4	1	

2975	267	2	0
3903	203	2	0
8437	146	4	1
...	...	...	...
1229	148	3	0
10593	151	3	0
12248	228	5	0
3147	154	2	0
6623	184	3	0

	promotion_last_5years	sales_IT	sales_RandD	sales_accounting	...	\
6958	0	0	0	0	...	
7534	0	0	0	0	...	
2975	0	0	1	0	...	
3903	0	0	0	0	...	
8437	0	0	0	0	...	
...	...	...	...	...	...	
1229	0	0	0	0	...	
10593	0	1	0	0	...	
12248	0	0	0	0	...	
3147	0	0	0	0	...	
6623	0	1	0	0	...	

	sales_management	sales_marketing	sales_product_mng	sales_sales	\
6958	0	0	0	1	
7534	0	0	0	0	
2975	0	0	0	0	
3903	0	0	0	1	
8437	0	0	0	0	
...	...	...	...	...	
1229	0	0	0	1	
10593	0	0	0	0	
12248	0	0	0	1	
3147	0	0	0	1	
6623	0	0	0	0	

	sales_support	sales_technical	salary_high	salary_low	salary_medium	\
6958	0	0	1	0	0	
7534	1	0	0	1	0	
2975	0	0	0	1	0	
3903	0	0	0	1	0	
8437	0	1	0	1	0	
...	...	...	...	...	...	
1229	0	0	0	1	0	
10593	0	0	1	0	0	
12248	0	0	0	1	0	
3147	0	0	0	0	1	

6623	0	0	0	1	0
------	---	---	---	---	---

	leave
6958	0
7534	0
2975	0
3903	0
8437	0
...	...
1229	1
10593	0
12248	1
3147	0
6623	0

[3000 rows x 21 columns]

```
[121]: # Based on the below probability score range, categorize the employees into
      ↪ four zones
      # suggest your thoughts on the retention strategies for each zone
```

```
zone=[]
prob=[]

for i in predicted_prob[:,1]:
    prob.append(i)
    if (i<=0.2):
        zone.append("Safe Zone")
    elif (i>0.2 and i<=0.6):
        zone.append("Low Risk Zone")
    elif (i>0.6 and i<=0.9):
        zone.append("Medium Risk Zone ")
    else:
        zone.append("High Risk Zone ")
```


```
[122]: categories = ["Safe Zone","Low Risk Zone","Medium Risk Zone ","High Risk Zone "]
      color = ["Green","Yellow","Orange","Red"]
```

```
[123]: colordict = dict(zip(categories, color))

      clr = pd.DataFrame({"zone":zone,"probability":prob})
      clr["Color"] = clr["zone"].apply(lambda x: colordict[x])
      clr['zone'] = clr['zone'].astype(str)
```

```
[131]: color= clr["Color"].tolist()
      c = ["Green","Red","Orange","Yellow"]
```

```
plt.figure(figsize=(7,7))
sns.countplot(x=zone, palette=c)
plt.show()
```



[ ]: